

Generative Models

COMS 4995 - Neural Network and Deep Learning

Overview

There seems to be 3 types of generative models prevalent nowadays :

- GANs
- Reversible Architecture
- VAE & Diffusion models

But what is the idea that is **common** in all of them?

There is some z vector that is drawn from some known distribution like normal distribution $z \sim \mathcal{N}(0,1)$ and that is used to map the data distribution represented as p_x

This mapping $p_z \rightarrow p_x$ is done by some function f_θ that is usually a neural network parameterized by θ .

So you end up with $f_\theta(z) = x'$ that forms a distribution $p_{x'}$

What is the loss here ? You want the distribution $p_{x'}$ to be as close to the true distribution of the data p_x . So we measure the loss as the discrepancy or how much does $p_{x'}$ differ from p_x

GANs

There is something known as Generator, $G(z)$ that takes the input $z \sim \mathcal{N}(0,1)$ and outputs a vector, x' that is supposed to mimic the actual distribution.

How do we measure if this value x' is close to a sample drawn from the actual distribution $x \sim p_x$

Therefore we have a discriminator, $D(x, x')$ that takes on the actual sample, x and the 'fake' sample x' and outputs a score to each of them telling which one seems to be from the actual distribution.

There various loss that has been defined for estimating $D(x, x')$:

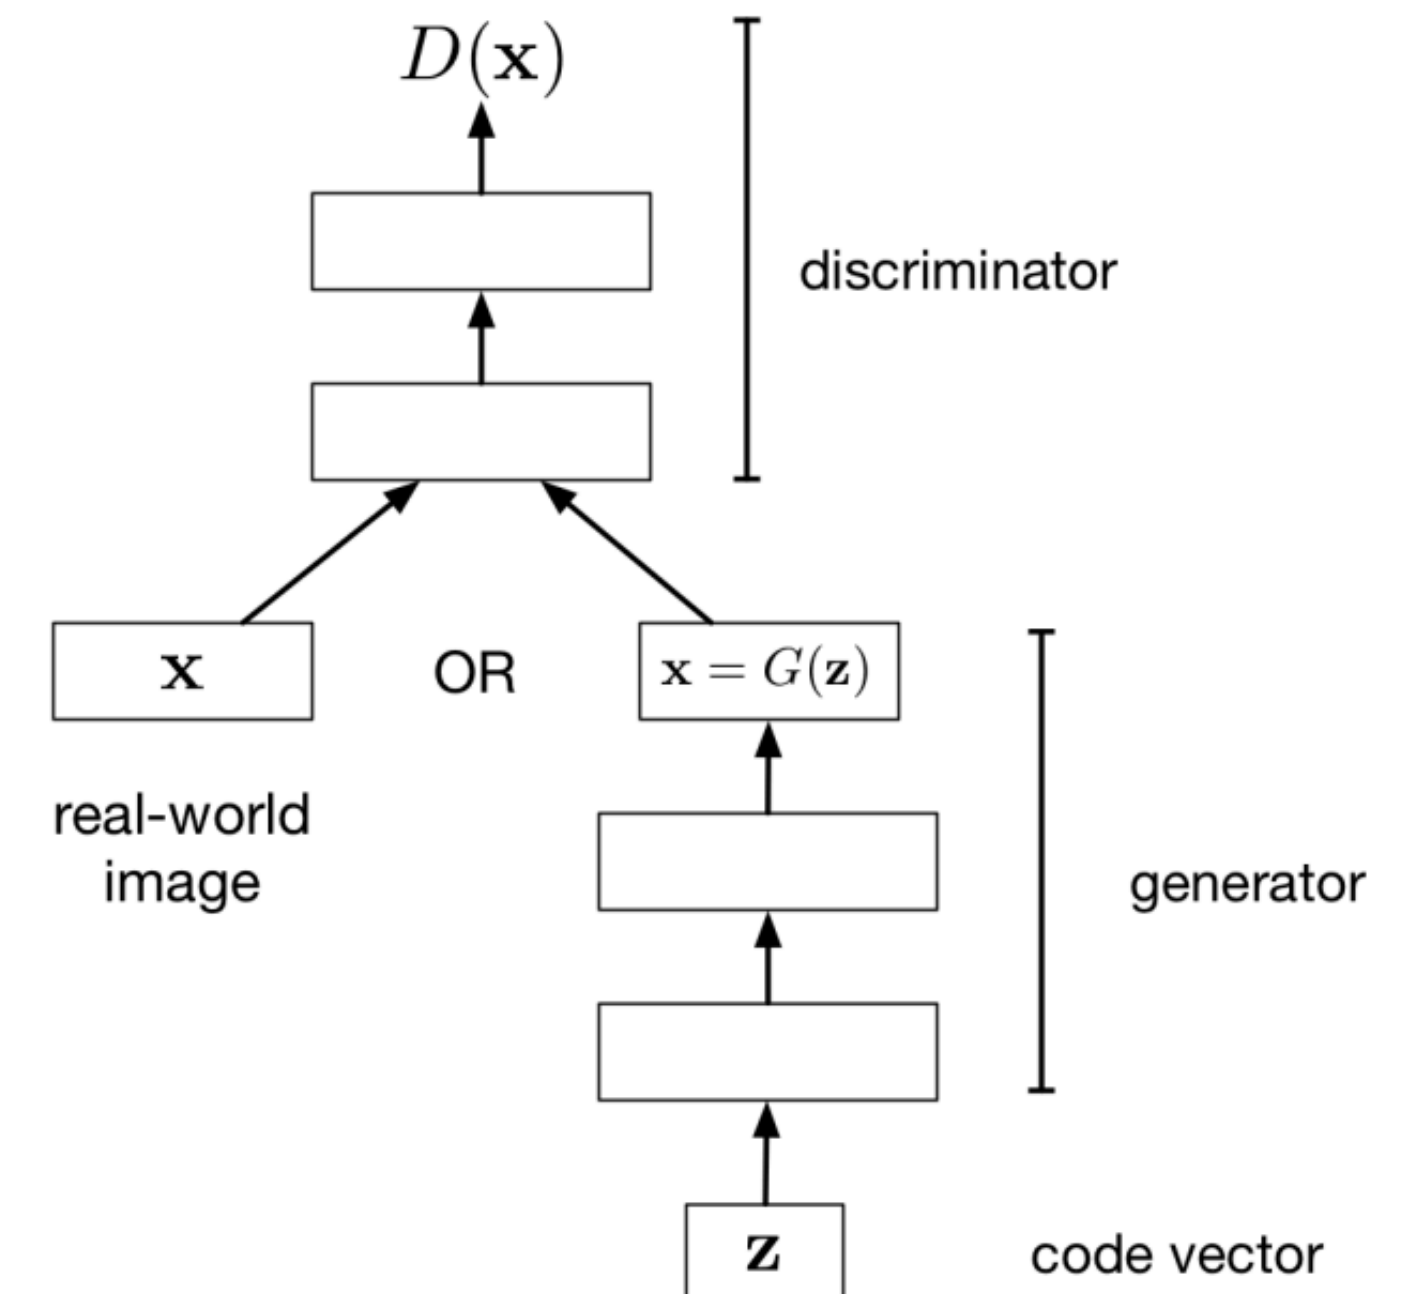
cross_entropy_loss

$$\mathcal{J}_D = \mathbf{E}_{x \sim p_x}[-\log(D(x))] + \mathbf{E}_z[-\log(1 - D(G(z)))]$$

$$\mathcal{J}_G = -\mathcal{J}_D$$

And then we use this to formulate the minimax problem


$$\max_G \min_D \mathcal{J}_D$$



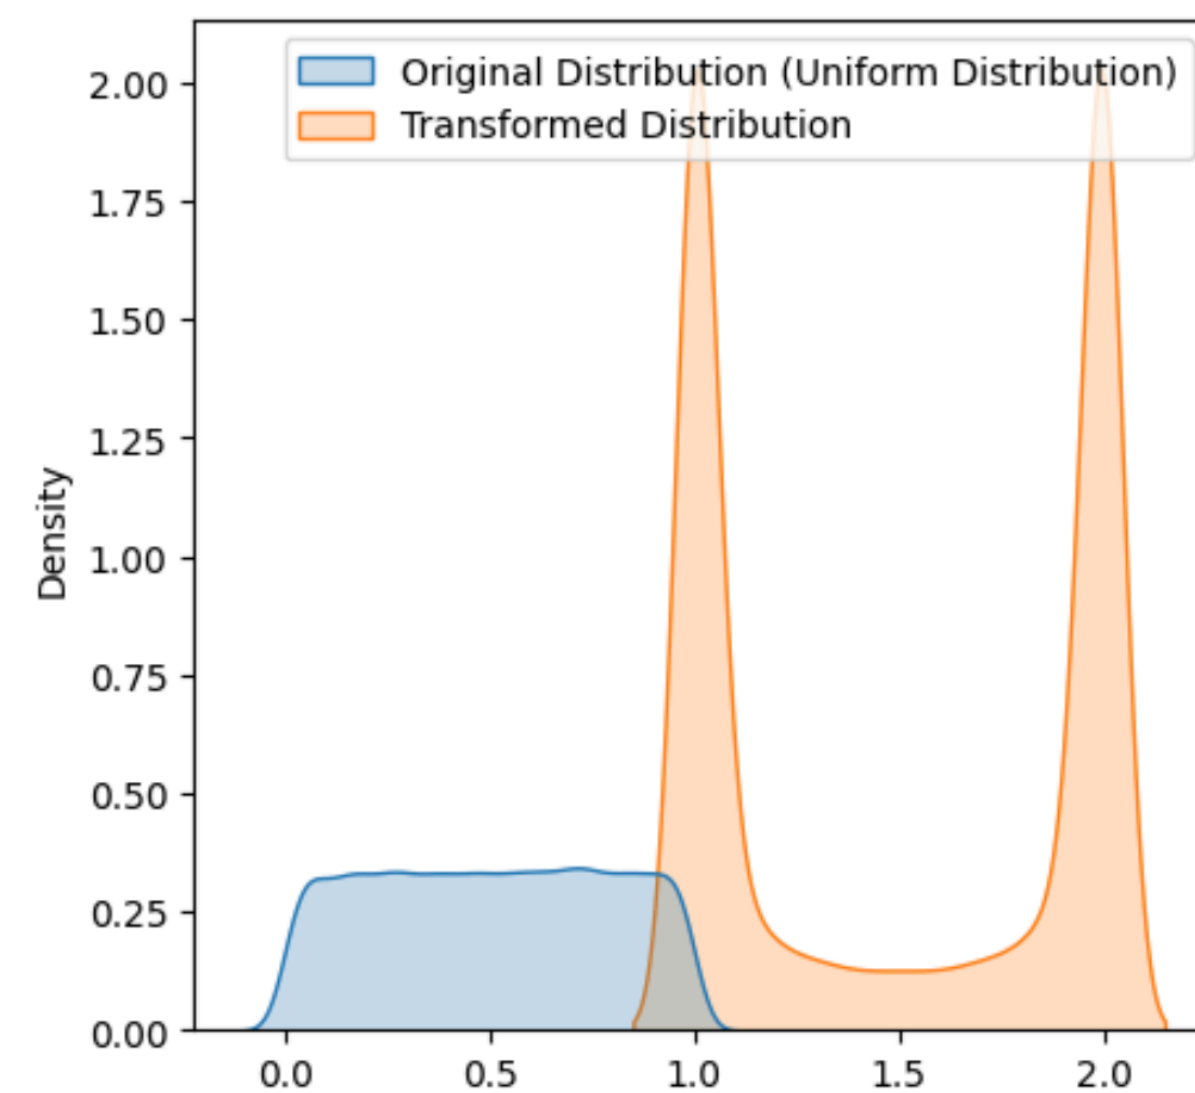
Some Examples and Codes

This things really work?

Transformed Distribution : p_x



Single Mode



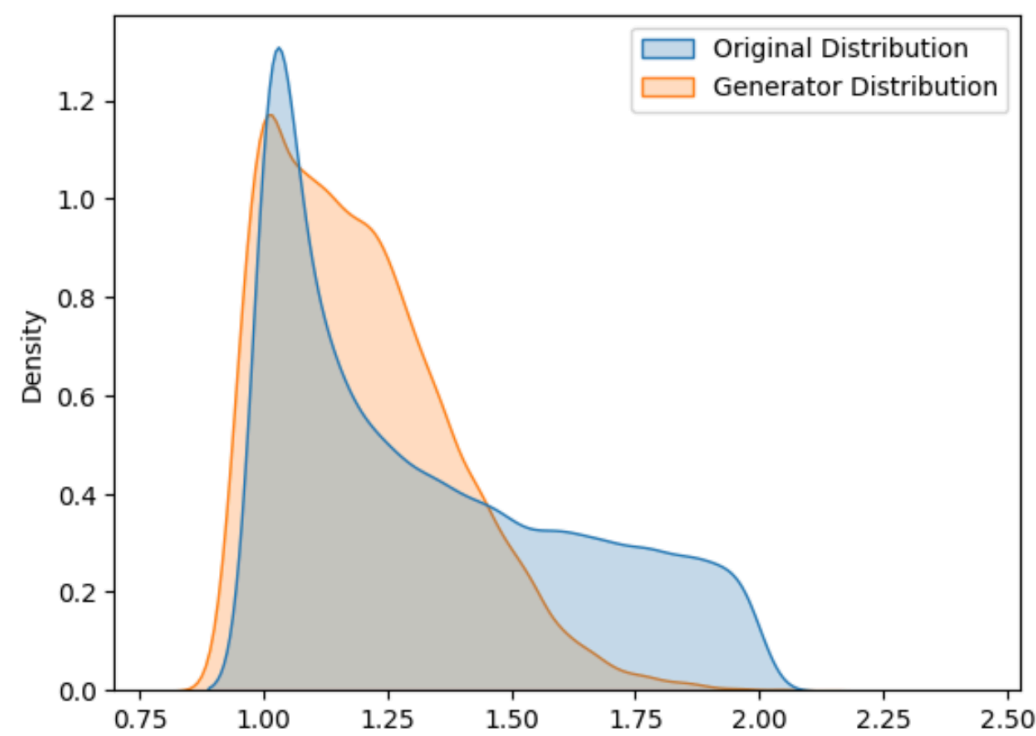
Bimodal

Recall : What are we trying to learn ????

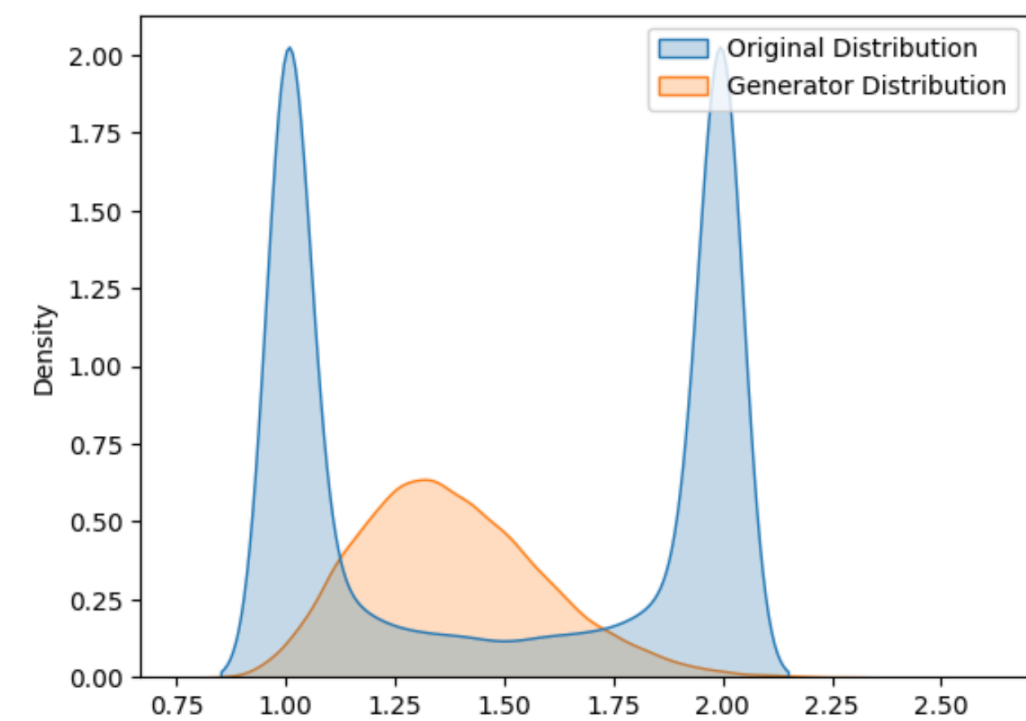
A function that takes in a value $z \sim N(0,1)$
and gives a sample x' that is close to the true
distribution p_x

Loss Function and Mode Collapse

Did Cross Entropy Loss work in all cases ?



Yeah probably for single mode



Not quite for bimodal


Mode Collapse !!

Wasserstein Loss


The output of the discriminator is a score

Discriminator loss : $D(x) - D(G(z))$

Generator loss : $D(G(z))$



Earth Mover distance



What all to be aware of ?

Do not loss your sleep over this

Training GAN is hard. Its an understatement. Its really hard to get a stable training

The **learning rate**, yes even the **epsilon values of Adam** needs to be tuned for this

Loss function, like **Wasserstein Loss** can help preventing in mode collapse

Unrolling of GAN : Can help in speed up

Clip the gradients at each updates

Normalizing Flow

Also called as reversible models. X is broken into $[x_1, x_2]$

$$p_X(x) = p_Z(f(x)) \left| \det\left(\frac{\partial f(x)}{\partial x^T}\right) \right|$$

In this you have something called as forward propagation which is the f that maps $x \rightarrow z$ that is being learnt.

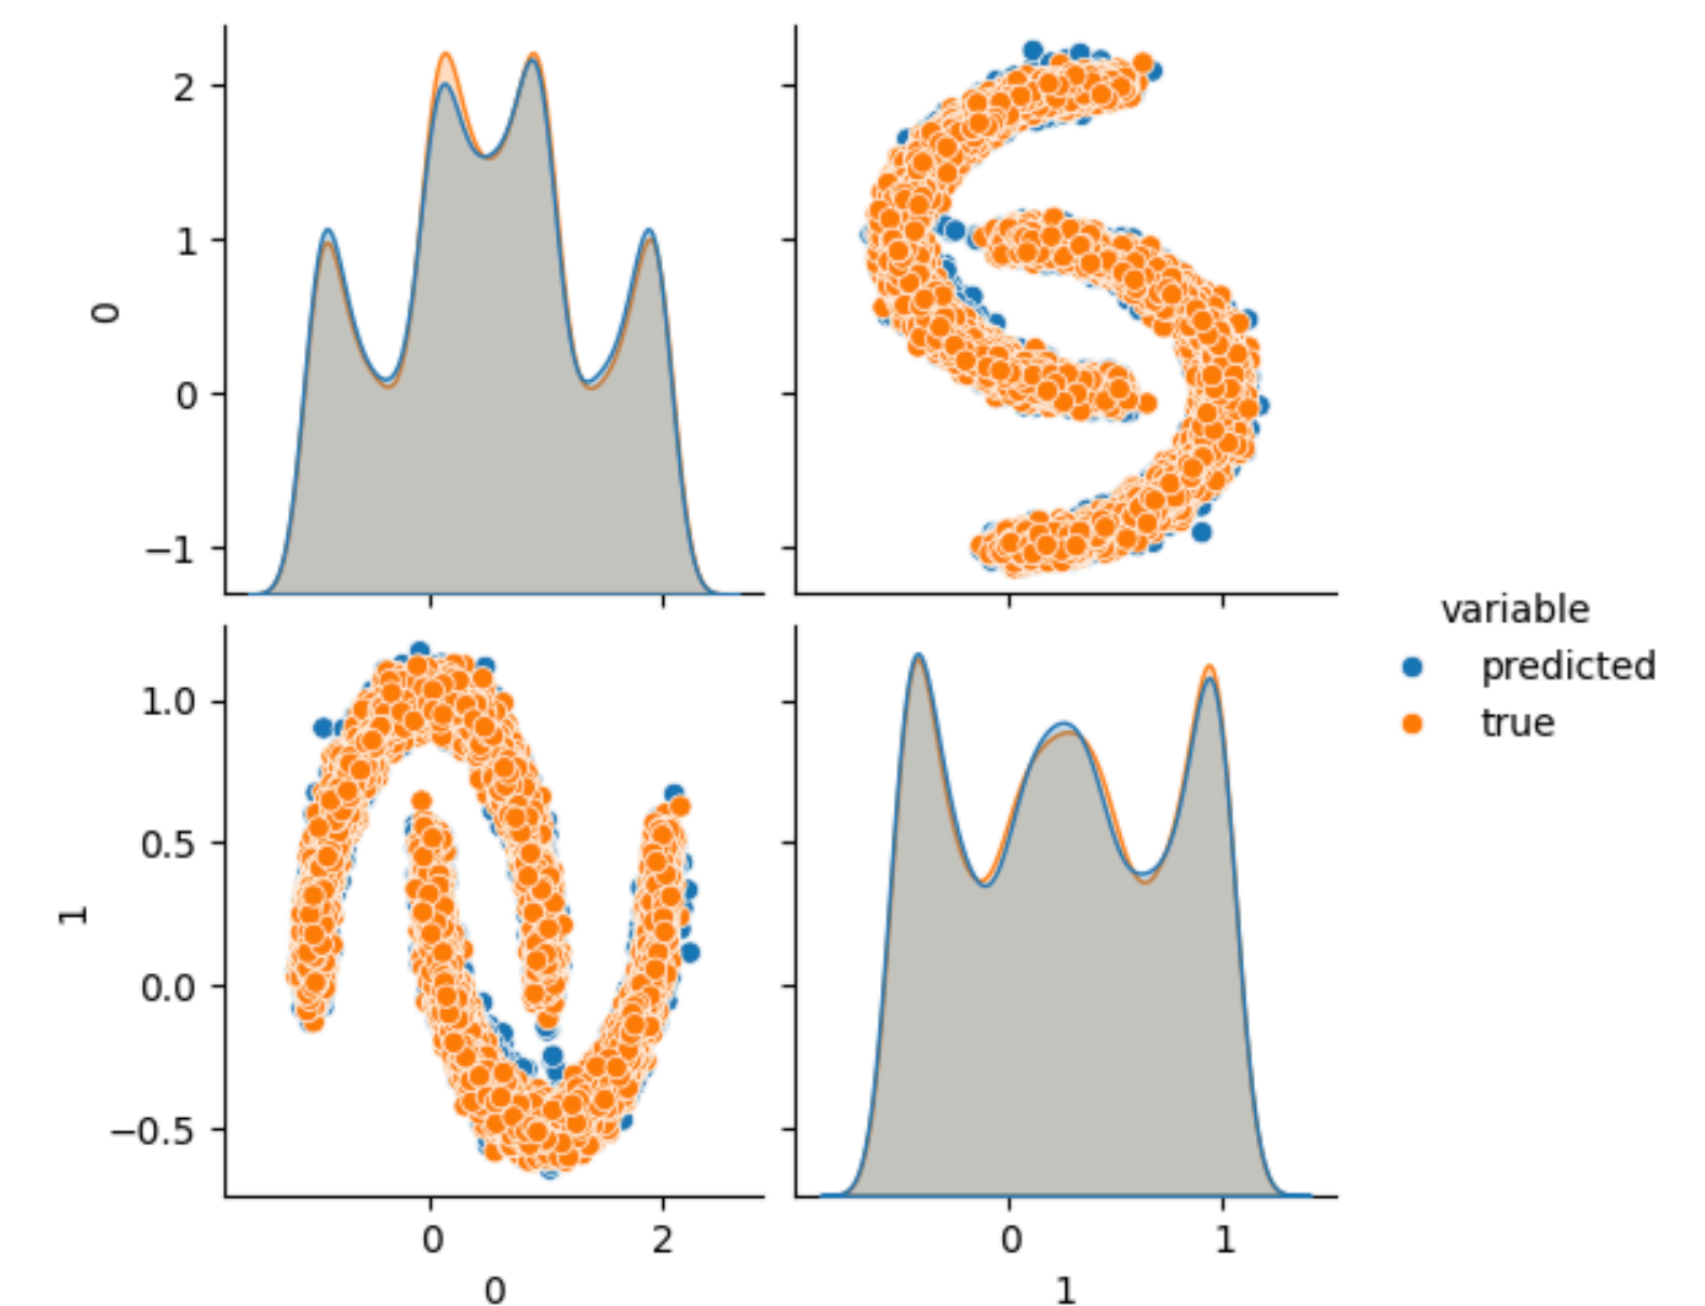

This f is **invertible**. So how do design **neural networks that are invertible** ??

So now our basic neural network is learning the transformations s and t .

These transformations forms one block. We can now have one more block over this and so on

In order to generate new samples from $z \sim N(0,1)$, we run through the inverse of these functions through the inverse propagation.

This forces each block to be of the same dimension as the input dimension



Variational Auto-Encoder

The overall idea is quite similar.

We take a sample $x \sim p_x$ and then get a representation of it z . However, we do not directly predict z , rather the distribution parameters from which z is sampled from μ, σ .

Then we use the re-parameterization trick to get a sample z :

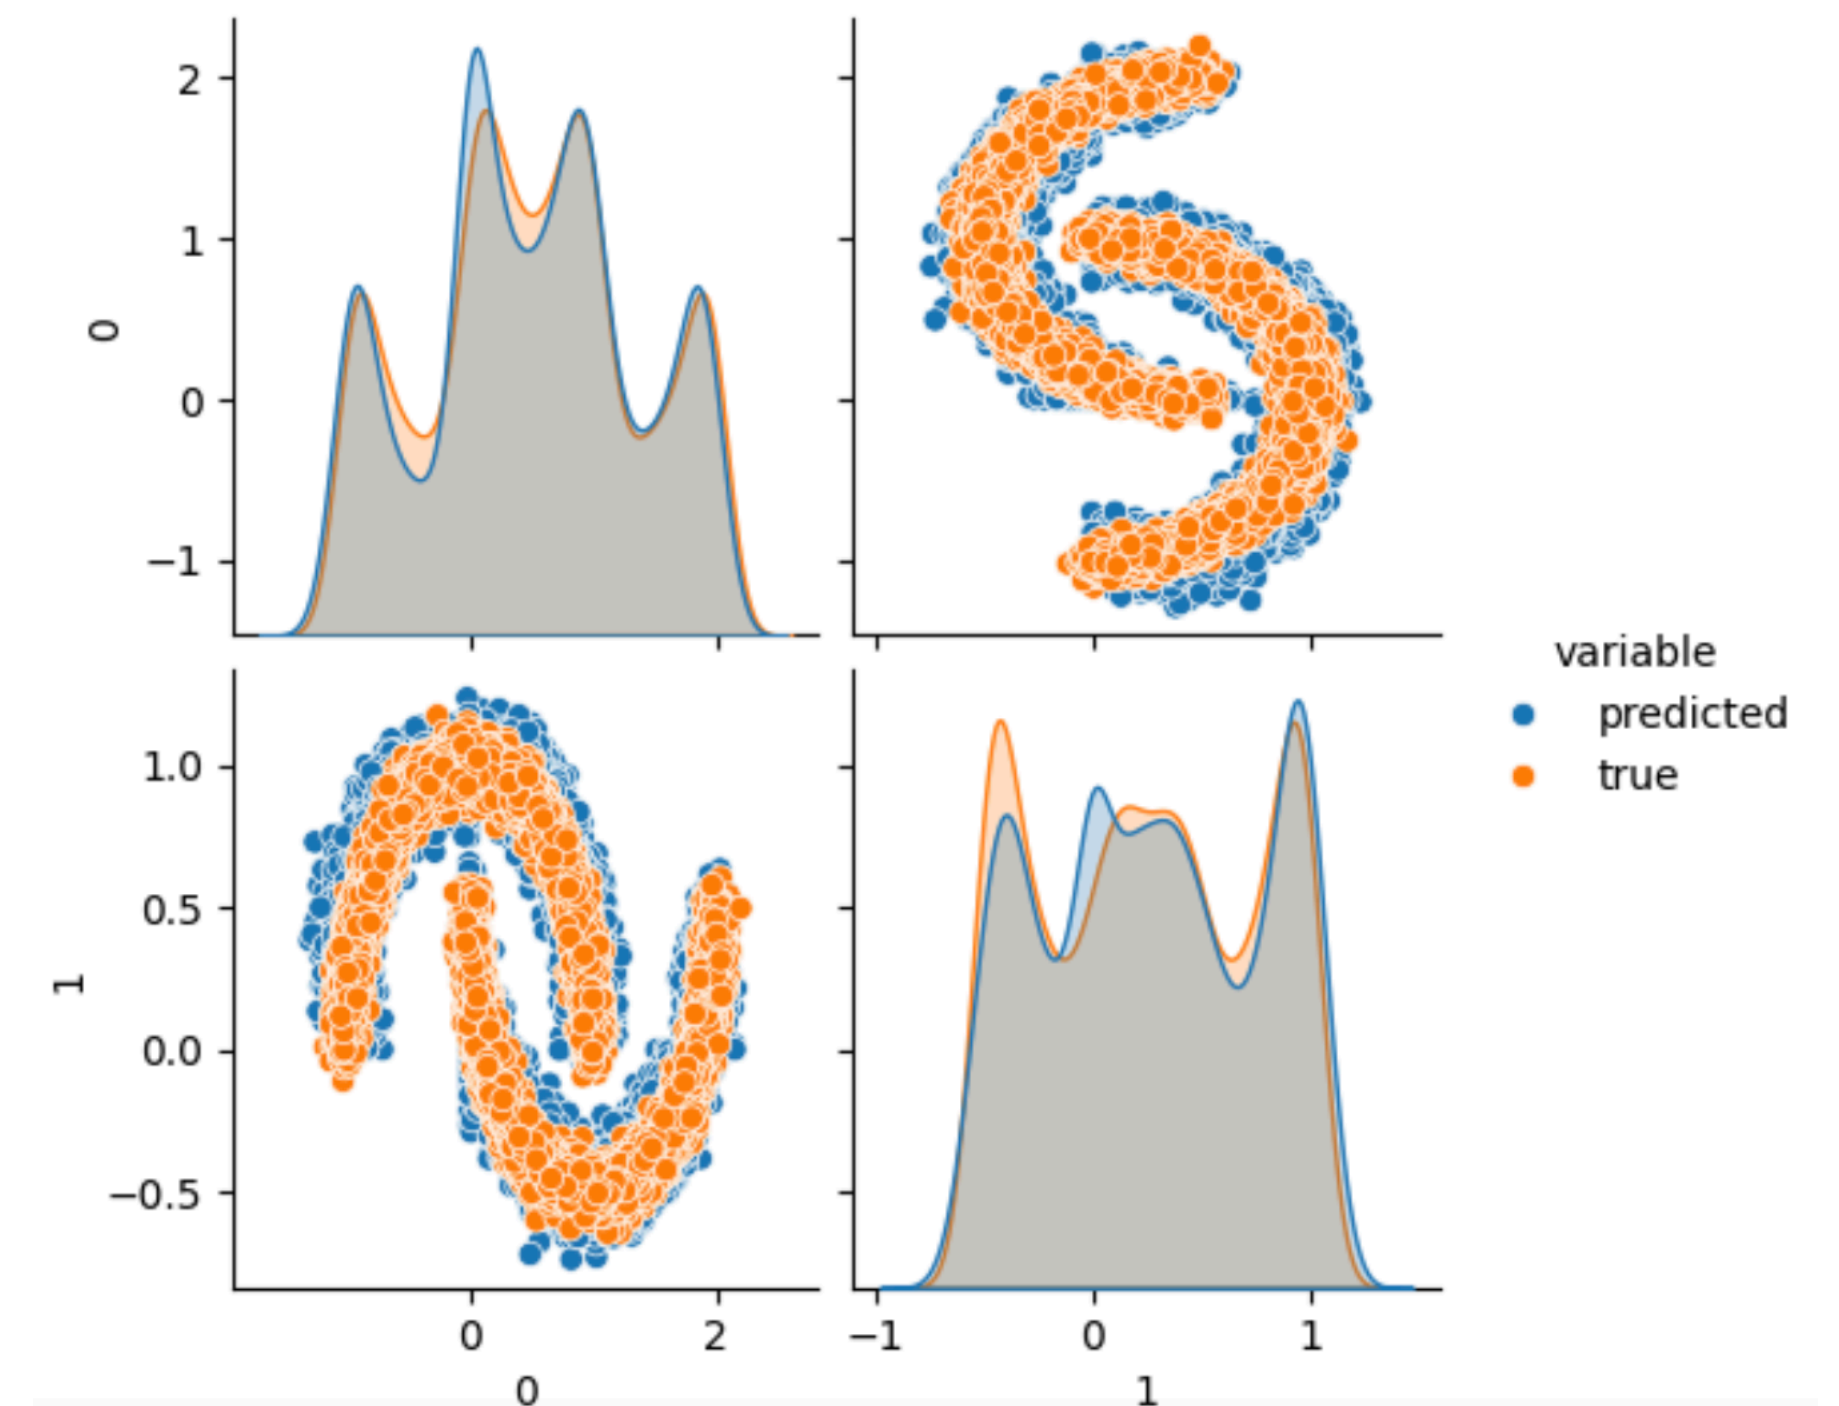
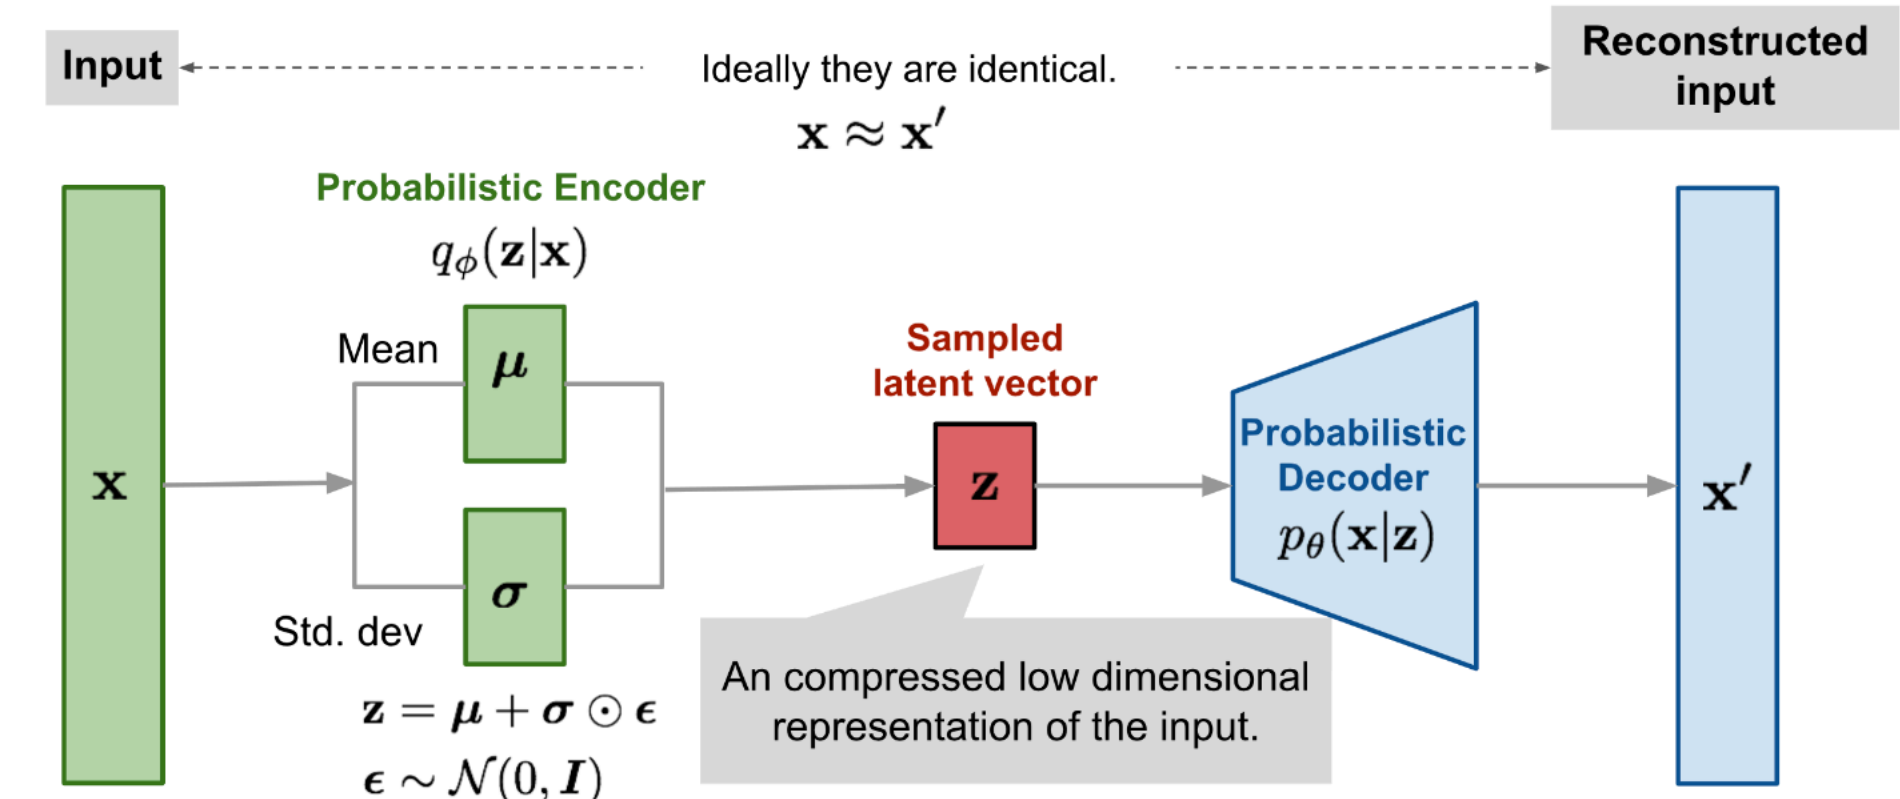
$$z = \mu + \sigma \odot \epsilon, \epsilon \sim N(0,1)$$

Then we pass through the decoder to produce a sample x' .

$$\mathbf{E}_q[\log \frac{p(z)}{q(z)}] + \mathbf{E}_q[\log p(x|z)]$$

KL Divergence Loss This is the MSE loss

The derivation of this itself is a maths problem Reconstruction Loss



Diffusion Model

We sample a point $x \sim p_x$ from the real distribution. A forward diffusion process involves adding a small amount of gaussian noise to the sample at each step controlled by variance scheduler β_t . Eventually when $T \rightarrow \infty$, x_T resembles a isotropic gaussian ($\Sigma = \sigma^2 I$)

$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I}) \quad q(\mathbf{x}_{1:T}|\mathbf{x}_0) = \prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1})$$

So now, we can do the reverse of this as well, called as the reverse diffusion process which takes a isotropic gaussian sample which becomes your code vector z and then gradually form the image or the sample which belongs to the real distribution.

There is a way to formulate x_t in terms of x_0 where $\alpha = 1 - \beta$

$$\begin{aligned} \mathbf{x}_t &= \sqrt{\alpha_t}\mathbf{x}_{t-1} + \sqrt{1 - \alpha_t}\boldsymbol{\epsilon}_{t-1} \\ &= \sqrt{\alpha_t\alpha_{t-1}}\mathbf{x}_{t-2} + \sqrt{1 - \alpha_t\alpha_{t-1}}\bar{\boldsymbol{\epsilon}}_{t-2} \\ &= \dots \\ &= \sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\boldsymbol{\epsilon} \end{aligned}$$

where $\boldsymbol{\epsilon}_{t-1}, \boldsymbol{\epsilon}_{t-2}, \dots \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

where $\bar{\boldsymbol{\epsilon}}_{t-2}$ merges two Gaussians (*).

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})$$

Post this the maths becomes too complicated.

Just follow this algorithm

Algorithm 1 Training	Algorithm 2 Sampling
<div>1: repeat 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 3: $t \sim \text{Uniform}(\{1, \dots, T\})$ 4: $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 5: Take gradient descent step on $\nabla_{\theta} \left\ \boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\boldsymbol{\epsilon}, t) \right\ ^2$ 6: until converged</div>	<div>1: $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 2: for $t = T, \dots, 1$ do 3: $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$ 4: $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 5: end for 6: return \mathbf{x}_0</div>